

AI Governance Policy Review Report

Organisation: BrightPath Skills Training
Build 6 · BrightPath ChatGPT Mastery Project

Date: 2026-06-17

Prototype status: production-style AI governance review prototype — not a production compliance, legal, safeguarding, HR, data-protection, or professional advisory system.

Synthetic scenarios only. Human review required before any real-world use.

Report Analytics Summary

Export Completion Status

Output	Status
Policy Pack	Available
Governance Framework	Available
Coverage Review	Available
Gap Analysis	Available
Recommendations	Available
Governance Maturity	Available
Governance Report	Available
Completion	100.0%

Policy Coverage Level Distribution

Coverage Level	Domains
Strong coverage	9
Partial coverage	3
Weak coverage	0
Not covered	0

Policy Gap Severity Distribution

Gap Severity	Count
Critical gap	0
High gap	0
Medium gap	2
Low gap	1

Recommendation Priority Distribution




Priority	Count
Urgent	0
High priority	0
Medium priority	3
Low priority	0

Governance Score Overview




Metric	Value
Coverage Score	83.2/100
Governance Score	79.4/100
Maturity Level	Managed governance

Report Quality: 10 of 10 sections present (100.0%)

Governance Analytics Charts



Policy Gap Severity Distribution



Full Governance Report

AI Governance Policy Review Report

Organisation: BrightPath Skills Training

Date: 2026-06-17

Cover Page

AI Governance Policy Review Report

Organisation: BrightPath Skills Training

Organisation type: Small UK training provider

Sector: Education and skills training

Country context: United Kingdom

Policy pack: Synthetic Responsible AI Policy Pack

Date generated: 2026-06-17

Prototype status: Production-style AI governance review prototype, not a production compliance, legal, safeguarding, HR, data-protection, or professional advisory system.

This report is generated from synthetic/demo policy text only. Human review is required before any real-world use.

Table of Contents

- 1. Executive Summary
 - 2. Policy Pack Overview
 - 3. Responsible AI Governance Framework
 - 4. Policy Coverage Review
 - 5. Policy Gap Analysis
 - 6. Policy Improvement Recommendations
 - 7. Governance Score and Maturity Summary
 - 8. Recommended Next Steps
 - 9. Responsible-Use Boundaries
 - 10. Prototype Limitations
 - 11. Appendices
-

1. Executive Summary

This report presents the findings of an AI governance policy review for BrightPath Skills Training. The review assessed the organisation's AI policy pack against a responsible AI governance framework covering 12 domains: strategy and ownership, approved AI tools, prohibited uses, data protection and confidentiality, learner and client data boundaries, safeguarding boundaries, human review and accountability, accuracy and hallucination control, bias, fairness and inclusion, staff training and capability, escalation and incident reporting, and monitoring and continuous improvement.

The overall policy coverage score is 83.2/100 (Strong coverage). This is based on deterministic keyword-based evidence matching across the synthetic policy pack — not a legal or regulatory assessment.

The policy pack contains useful foundations, but several areas need clearer operational detail before scaling AI use.

Key gap themes include: Data protection and confidentiality, Learner and client data boundaries, Bias, fairness, and inclusion, AI governance strategy and ownership.

The overall AI governance maturity level is Managed governance (score: 79.4/100). The organisation appears ready for structured pilots and limited scaling, subject to monitoring, staff training, and periodic governance review.

All outputs from this review are generated from synthetic/demo policy text only. Human review and responsible-owner approval are required before any policy change, governance decision, or adoption guidance is acted upon.

2. Policy Pack Overview

Organisation: BrightPath Skills Training
Organisation type: Small UK training provider
Sector: Education and skills training
Country context: United Kingdom
Policy pack: Synthetic Responsible AI Policy Pack

Policies Reviewed

Policy	Type	Owner
AI Acceptable Use Policy	Acceptable Use	Chief Executive Officer
Data Protection and AI Use Guidance	Data Protection	Data Protection Lead
Safeguarding and AI Boundary Policy	Safeguarding	Designated Safeguarding Lead
Staff AI Tool Use Procedure	Procedure	Head of Operations
AI Output Review Checklist	Checklist	Quality and Standards Lead
AI Incident and Escalation Procedure	Incident Procedure	Head of Operations

Risk Areas Covered

- accountability
- accuracy and hallucination
- approved tools
- bias and fairness
- confidential data
- data protection
- escalation routes
- human review
- incident reporting
- learner/client data boundaries
- monitoring and review
- personal data

- prohibited uses
- safeguarding boundaries
- staff training
- supplier/tool approval

Synthetic data only. This policy pack is demo content for portfolio demonstration. Not an approved organisational policy.

3. Responsible AI Governance Framework

The responsible AI governance framework covers 12 domains, of which 7 are high priority and 5 are medium priority.

Domain ID	Domain	Priority
GOV-001	Strategy and Ownership	High
GOV-002	Approved AI Tools	High
GOV-003	Prohibited AI Uses	Medium
GOV-004	Data Protection and Confidentiality	High
GOV-005	Learner and Client Data Boundaries	High
GOV-006	Safeguarding Boundaries	High
GOV-007	Human Review and Accountability	High
GOV-008	Accuracy and Hallucination Control	Medium
GOV-009	Bias, Fairness, and Inclusion	Medium
GOV-010	Staff Training and Capability	Medium
GOV-011	Escalation and Incident Reporting	High
GOV-012	Monitoring, Review, and Continuous Improvement	Medium

The framework defines what good responsible AI governance looks like. Each domain specifies expected policy evidence and example controls that the policy coverage checker uses to assess the synthetic policy pack.

4. Policy Coverage Review

The policy coverage review assessed BrightPath Skills Training's AI policy pack against the 12 responsible AI governance domains using deterministic keyword-based evidence matching. The overall coverage score is 83.2/100 (Strong coverage) across 12 domains.

Overall coverage score: 83.2/100

Overall coverage level: Strong coverage

Domains checked: 12

Coverage Level	Count
Strong coverage	9
Partial coverage	3
Weak coverage	0
Not covered	0

Recommended focus areas:

- Review partial coverage areas for clarity and completeness.

Domain Coverage Results

Domain	Priority	Score	Level
Strategy and Ownership	High	60.0/100	Partial coverage
Approved AI Tools	High	84.5/100	Strong coverage
Prohibited AI Uses	Medium	88.7/100	Strong coverage
Data Protection and Confidentiality	High	92.8/100	Strong coverage
Learner and Client Data Boundaries	High	60.0/100	Partial coverage
Safeguarding Boundaries	High	92.8/100	Strong coverage
Human Review and Accountability	High	88.7/100	Strong coverage
Accuracy and Hallucination Control	Medium	88.7/100	Strong coverage
Bias, Fairness, and Inclusion	Medium	60.0/100	Partial coverage
Staff Training and Capability	Medium	97.0/100	Strong coverage
Escalation and Incident Reporting	High	92.8/100	Strong coverage
Monitoring, Review, and Continuous Improvement	Medium	92.8/100	Strong coverage

Coverage scores are based on deterministic keyword matching in synthetic policy text. They are not legal or regulatory assessments.

5. Policy Gap Analysis

The gap analysis reviewed BrightPath Skills Training's AI policy pack against the 12 responsible AI governance domains, identifying areas where policy coverage is absent, weak, or insufficient. Gaps are classified by severity and domain priority to support the consulting conversation about where to focus the organisation's policy review effort.

Total gaps identified: 3

Critical gaps: 0

High gaps: 0

Medium gaps: 2

Low gaps: 1

Domains with sufficient coverage: 9

The policy pack contains useful foundations, but several areas need clearer operational detail before scaling AI use.

Gap themes:

- Data protection and confidentiality
- Learner and client data boundaries
- Bias, fairness, and inclusion
- AI governance strategy and ownership

Recommended focus areas:

- Clarify learner/client data boundaries.

Prioritised Gaps

Gap ID	Domain	Severity	Priority	Score
GAP-001	Strategy and Ownership	Medium gap	High	60.0/100
GAP-002	Learner and Client Data Boundaries	Medium gap	High	60.0/100
GAP-003	Bias, Fairness, and Inclusion	Low gap	Medium	60.0/100

Gap severity classifications are indicative only. A Critical gap does not mean the organisation is non-compliant. Qualified human review is required.

6. Policy Improvement Recommendations

This section sets out policy improvement recommendations for BrightPath Skills Training, derived from the gap analysis findings. Each recommendation identifies a target policy, suggested responsible owner, and a wording direction to guide the policy review. Recommendations are prioritised by gap severity and domain priority so that the most important actions are visible at a glance.

Total recommendations: 3

Urgent: 0

High priority: 0

Medium priority: 3

Low priority: 0

Quick wins: 3

The policy pack has useful foundations, but operational details should be strengthened before AI use expands.

Policy update themes:

- Data protection and confidentiality
- Learner and client data boundaries
- Bias and fairness review
- AI governance ownership

Key responsible owners:

- Chief Executive Officer / Senior Leadership
- Data Protection Lead / Operations Manager
- Equality Lead / Quality Lead / Senior Manager

Quick Wins

- REC-001 — Review and clarify existing wording: Strategy and Ownership (Medium priority)
- REC-002 — Review and clarify existing wording: Learner and Client Data Boundaries (Medium priority)
- REC-003 — Review and clarify existing wording: Bias, Fairness, and Inclusion (Medium priority)

Prioritised Recommendations

ID	Domain	Priority	Action	Owner
REC-001	Strategy and Ownership	Medium priority	Review and clarify existing policies	Chief Executive Officer / Senior Leadership
REC-002	Learner and Client Data Boundaries	Medium priority	Review and clarify existing policies	Data Protection Lead / Operations Manager
REC-003	Bias, Fairness, and Inclusion	Medium priority	Review and clarify existing policies	Equality Lead / Quality Lead / Senior Manager

Wording directions are not legally approved policy text. They are starting points for review by appropriate responsible owners.
Human review is required before any policy change.

7. Governance Score and Maturity Summary

This section presents the AI governance maturity assessment for BrightPath Skills Training. The maturity score is derived from policy coverage scores, with penalties applied for identified gap severity and recommendation priority, giving a calibrated picture of governance readiness across the 12 assessed domains.

Overall governance score: 79.4/100

Maturity level: Managed governance

The organisation has strong governance coverage across most reviewed areas. AI pilots and limited scaling may be appropriate with monitoring, training, incident reporting, and periodic review.

Adoption readiness position:

The organisation appears ready for structured pilots and limited scaling, subject to monitoring, staff training, and periodic governance review.

Recommended next step:

Evidence monitoring cycles and continuous improvement reviews to consolidate Managed governance maturity.

Governance strengths:

- Staff Training and Capability — 97/100 (Optimised governance)
- Data Protection and Confidentiality — 92/100 (Optimised governance)
- Safeguarding Boundaries — 92/100 (Optimised governance)
- Escalation and Incident Reporting — 92/100 (Optimised governance)
- Monitoring, Review, and Continuous Improvement — 92/100 (Optimised governance)
- Prohibited AI Uses — 88/100 (Managed governance)
- Human Review and Accountability — 88/100 (Managed governance)
- Accuracy and Hallucination Control — 88/100 (Managed governance)
- Approved AI Tools — 84/100 (Managed governance)

Governance weaknesses:

- Strategy and Ownership — 47/100 (Developing governance)
- Learner and Client Data Boundaries — 47/100 (Developing governance)

Domain Maturity Scores

Domain	Priority	Coverage	Maturity Score	Maturity Level
Strategy and Ownership	High	60/100	47/100	Developing governance
Approved AI Tools	High	84/100	84/100	Managed governance

Domain	Priority	Coverage	Maturity Score	Maturity Level
Prohibited AI Uses	Medium	88/100	88/100	Managed governance
Data Protection and Confidentiality	High	92/100	92/100	Optimised governance
Learner and Client Data Boundaries	High	60/100	47/100	Developing governance
Safeguarding Boundaries	High	92/100	92/100	Optimised governance
Human Review and Accountability	High	88/100	88/100	Managed governance
Accuracy and Hallucination Control	Medium	88/100	88/100	Managed governance
Bias, Fairness, and Inclusion	Medium	60/100	52/100	Defined governance
Staff Training and Capability	Medium	97/100	97/100	Optimised governance
Escalation and Incident Reporting	High	92/100	92/100	Optimised governance
Monitoring, Review, and Continuous Improvement	Medium	92/100	92/100	Optimised governance

Maturity levels are indicative score bands, not formal compliance ratings or regulatory assessments. Human review is required.

8. Recommended Next Steps

1. Evidence monitoring cycles and continuous improvement reviews to consolidate Managed governance maturity.
2. Strengthen Strategy and Ownership policy coverage (current maturity score: 47/100).
3. Strengthen Learner and Client Data Boundaries policy coverage (current maturity score: 47/100).
4. Review and evidence monitoring, training, and continuous improvement cycles to move towards Optimised governance.
5. Clarify the current wording and add missing operational details so staff know what to do in practice. Review this gap before scaling AI use or approving wider staff adoption.
6. Review the highest-priority gaps with responsible owners.
7. Clarify approved and prohibited AI use cases in the AI Acceptable Use Policy.
8. Strengthen learner and client data boundary wording in the Data Protection and AI Use Guidance.

All next steps should be reviewed and approved by qualified responsible owners before any real-world action is taken.

9. Responsible-Use Boundaries

This governance report is generated from synthetic/demo policy text only. It must not be used with real client policies, learner data, safeguarding case details, staff HR data, personal data, confidential data, or regulated information without appropriate governance, approvals, and responsible owners.

This prototype does not provide legal, safeguarding, HR, compliance, data-protection, financial, medical, academic-integrity, or professional governance advice. Suggested wording directions are not legally approved policy text. They require review by appropriate responsible owners before real-world use. This is a deterministic governance review support tool, not a compliance certification system. Human review remains required before any real-world use.

Suggested wording directions are not legally approved policy text. They require review by appropriate responsible owners before real-world use.

Human review remains required before any real-world use.

10. Prototype Limitations

This governance review prototype has the following limitations:

- Synthetic/demo policy data only. The policy pack is demo content, not real organisational policies.
- Deterministic keyword-based coverage review. Coverage scores are based on keyword frequency in synthetic text, not semantic, legal, or compliance analysis.
- No legal or regulatory analysis. Coverage scores, gap severities, maturity levels, and recommendations are not legal or regulatory risk assessments.
- No external AI or LLM API calls. All outputs are deterministic and template-based. No OpenAI, Claude, LangChain, LlamaIndex, or similar services are used.
- No real policy validation. This tool does not validate real-world policy adequacy, compliance, or legal sufficiency.
- No legal, compliance, safeguarding, HR, or data-protection approval. Outputs must be reviewed by qualified professionals.
- Not production deployed. This is a local portfolio prototype with no authentication, audit logging, or persistent storage.
- Human review required. All outputs require responsible human review before any real-world governance action.

11. Appendices

A. Policy List

- POL-001 — AI Acceptable Use Policy (Owner: Chief Executive Officer)
- POL-002 — Data Protection and AI Use Guidance (Owner: Data Protection Lead)
- POL-003 — Safeguarding and AI Boundary Policy (Owner: Designated Safeguarding Lead)
- POL-004 — Staff AI Tool Use Procedure (Owner: Head of Operations)
- POL-005 — AI Output Review Checklist (Owner: Quality and Standards Lead)
- POL-006 — AI Incident and Escalation Procedure (Owner: Head of Operations)

B. Governance Domain List

- GOV-001 — Strategy and Ownership (High priority)
- GOV-002 — Approved AI Tools (High priority)
- GOV-003 — Prohibited AI Uses (Medium priority)
- GOV-004 — Data Protection and Confidentiality (High priority)
- GOV-005 — Learner and Client Data Boundaries (High priority)
- GOV-006 — Safeguarding Boundaries (High priority)
- GOV-007 — Human Review and Accountability (High priority)
- GOV-008 — Accuracy and Hallucination Control (Medium priority)
- GOV-009 — Bias, Fairness, and Inclusion (Medium priority)
- GOV-010 — Staff Training and Capability (Medium priority)
- GOV-011 — Escalation and Incident Reporting (High priority)
- GOV-012 — Monitoring, Review, and Continuous Improvement (Medium priority)

C. Coverage Results

Domain	Score	Level
Strategy and Ownership	60.0/100	Partial coverage

Domain	Score	Level
Approved AI Tools	84.5/100	Strong coverage
Prohibited AI Uses	88.7/100	Strong coverage
Data Protection and Confidentiality	92.8/100	Strong coverage
Learner and Client Data Boundaries	60.0/100	Partial coverage
Safeguarding Boundaries	92.8/100	Strong coverage
Human Review and Accountability	88.7/100	Strong coverage
Accuracy and Hallucination Control	88.7/100	Strong coverage
Bias, Fairness, and Inclusion	60.0/100	Partial coverage
Staff Training and Capability	97.0/100	Strong coverage
Escalation and Incident Reporting	92.8/100	Strong coverage
Monitoring, Review, and Continuous Improvement	92.8/100	Strong coverage

D. Prioritised Gap Table

Gap ID	Domain	Severity	Score
GAP-001	Strategy and Ownership	Medium gap	60.0/100
GAP-002	Learner and Client Data Boundaries	Medium gap	60.0/100
GAP-003	Bias, Fairness, and Inclusion	Low gap	60.0/100

E. Recommendation Table

ID	Domain	Priority	Action
REC-001	Strategy and Ownership	Medium priority	Review and clarify existing wording
REC-002	Learner and Client Data Boundaries	Medium priority	Review and clarify existing wording
REC-003	Bias, Fairness, and Inclusion	Medium priority	Review and clarify existing wording

F. Domain Maturity Score Table

Domain	Maturity Score	Maturity Level
Strategy and Ownership	47/100	Developing governance
Approved AI Tools	84/100	Managed governance
Prohibited AI Uses	88/100	Managed governance
Data Protection and Confidentiality	92/100	Optimised governance
Learner and Client Data Boundaries	47/100	Developing governance
Safeguarding Boundaries	92/100	Optimised governance
Human Review and Accountability	88/100	Managed governance
Accuracy and Hallucination Control	88/100	Managed governance
Bias, Fairness, and Inclusion	52/100	Defined governance

Domain	Maturity Score	Maturity Level
Staff Training and Capability	97/100	Optimised governance
Escalation and Incident Reporting	92/100	Optimised governance
Monitoring, Review, and Continuous Improvement	92/100	Optimised governance

G. Source Outputs Availability

- ■ Synthetic policy pack
- ■ Responsible AI governance framework
- ■ Policy coverage review
- ■ Policy gap analysis
- ■ Policy improvement recommendations
- ■ Governance maturity summary

Build 6 · AI Governance Policy Checker · BrightPath ChatGPT Mastery Project

Synthetic scenarios only. Human review required before any real-world use.

Responsible-Use Boundaries

This governance report is generated from synthetic/demo policy text only. It must not be used with real client policies, learner data, safeguarding case details, staff HR data, personal data, confidential data, or regulated information without appropriate governance, approvals, and responsible owners. This prototype does not provide legal, safeguarding, HR, compliance, data-protection, financial, medical, academic-integrity, or professional governance advice. Suggested wording directions are not legally approved policy text. They require review by appropriate responsible owners before real-world use. This is a deterministic governance review support tool, not a compliance certification system. Human review remains required before any real-world use.

Build 6 · AI Governance Policy Checker · BrightPath ChatGPT Mastery Project

Synthetic scenarios only. Human review required before any real-world use.